

CS271: DATA STRUCTURES

Name: Bao Luu, Thanh Do, Khoa Nguyen

Instructor: Dr. Stacey Truex

Project #5

Exercises

Inserting and Deleting with B-Trees

- 24/26 1. Visually represent the result of inserting into a BTree with $t = 2$ the following values in the order listed:

$$F, J, D, H, B, C, I, G, A, E, K, L, M$$

For each value, show the BTree at the completion of the call to $\text{B-TREE-INSERT}(T, k)$. Additionally, for each value inserted, list next to the resulting B-Tree, the function calls made throughout the insertion. An example of this format is given at the end of this project description for your reference.

Insert F:

F		
---	--	--

$$\text{B-TREE-INSERT}(T, F)$$

$$\text{B-TREE-INSERT-NONFULL}(x = T.root, F)$$

Insert J:

F	J	
---	---	--

$$\text{B-TREE-INSERT}(T, J)$$

$$\text{B-TREE-INSERT-NONFULL}(x = T.root, J)$$

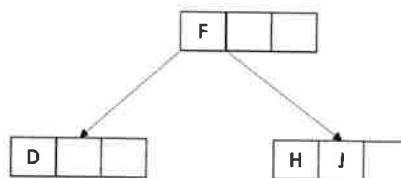
Insert D:

D	F	J
---	---	---

$$\text{B-TREE-INSERT}(T, D)$$

$$\text{B-TREE-INSERT-NONFULL}(x = T.root, D)$$

Insert H:

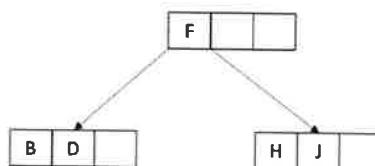


B-TREE-INSERT(T , H)

B-TREE-SPLIT-CHILD($x = T.root$, 1)

← B-TREE-INSERT-NONFULL(x , H) *this call is to insert into subtree rooted @ $s = T.root$ - will trigger an additional call to insert @ 2nd child*

Insert B:

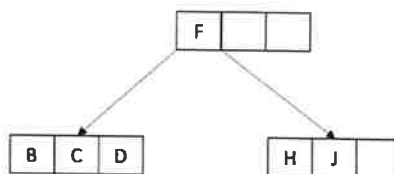


B-TREE-INSERT(T , B)

B-TREE-INSERT-NONFULL($x = T.root$, B)

B-TREE-INSERT-NONFULL($x = x.c_1$, B)

Insert C:

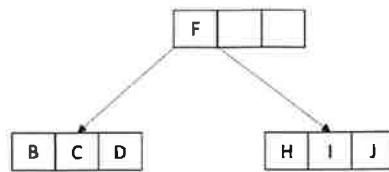


B-TREE-INSERT(T , C)

B-TREE-INSERT-NONFULL($x = T.root$, C)

B-TREE-INSERT-NONFULL($x = x.c_1$, C)

Insert I:

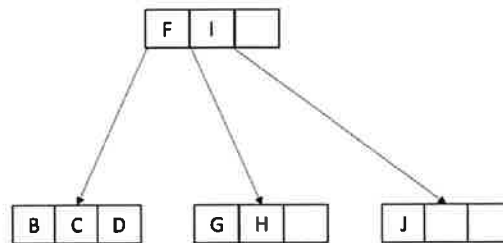


B-TREE-INSERT(T, I)

B-TREE-INSERT-NONFULL($x = T.root, I$)

B-TREE-INSERT-NONFULL($x = x.c_2, I$)

Insert G:



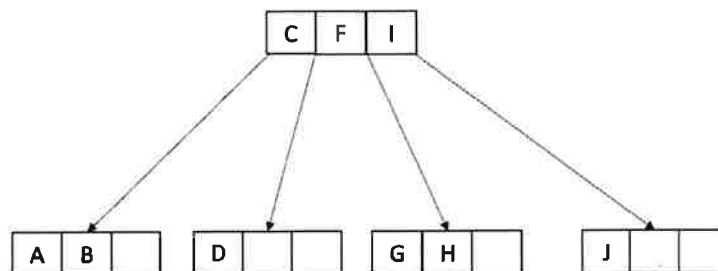
B-TREE-INSERT(T, G)

B-TREE-INSERT-NONFULL($x = T.root, G$)

B-TREE-SPLIT-CHILD($x, 2$)

B-TREE-INSERT-NONFULL($x = x.c_2, G$)

Insert A:



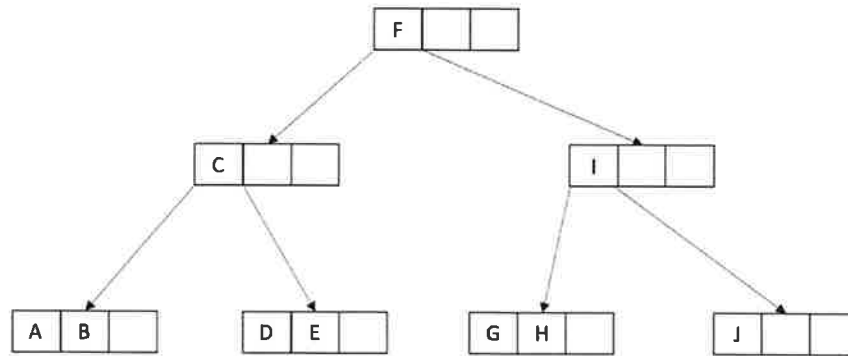
B-TREE-INSERT(T, A)

B-TREE-INSERT-NONFULL($x = T.root, A$)

B-TREE-SPLIT-CHILD($x, 1$)

B-TREE-INSERT-NONFULL($x = x.c_1, A$)

Insert E:



B-TREE-INSERT(T, E)

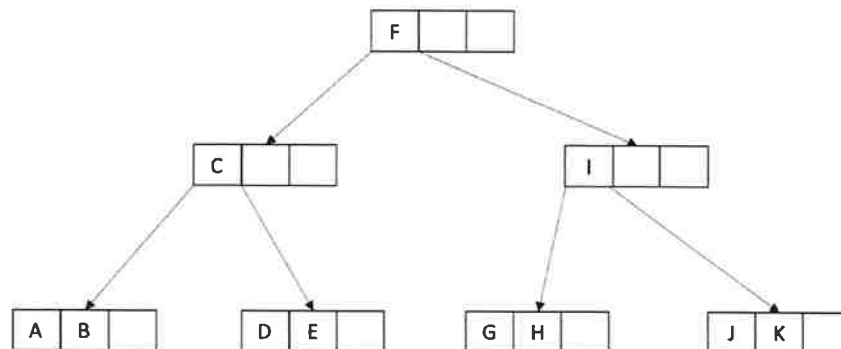
B-TREE-SPLIT-CHILD($x, 1$) *undefined in call to B-TREE-INSERT*

B-TREE-INSERT-NONFULL($x = T.root, E$)

B-TREE-INSERT-NONFULL($x = x.c_1, E$)

B-TREE-INSERT-NONFULL($x = x.c_2, E$)

Insert K:

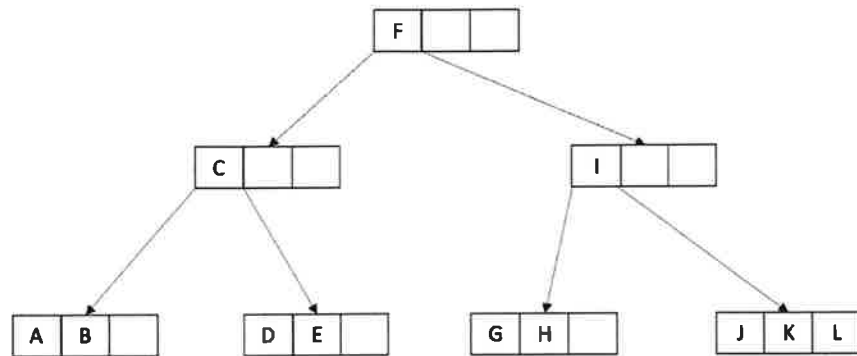


B-TREE-INSERT(T, K)

B-TREE-INSERT-NONFULL($x = T.root, K$)

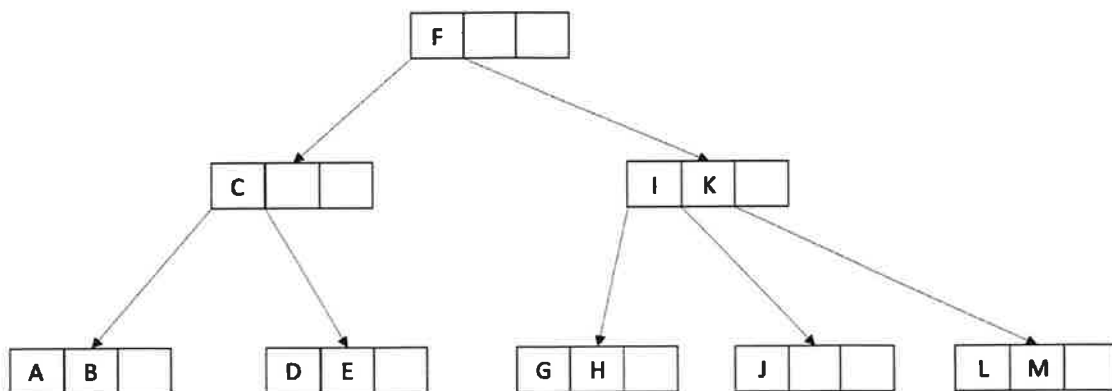
B-TREE-INSERT-NONFULL($x = x.c_2$, K)
 B-TREE-INSERT-NONFULL($x = x.c_2$, K)

Insert L:



B-TREE-INSERT(T , L)
 B-TREE-INSERT-NONFULL($x = T.root$, L)
 B-TREE-INSERT-NONFULL($x = x.c_2$, L)
 B-TREE-INSERT-NONFULL($x = x.c_2$, L)

Insert M:



B-TREE-INSERT(T , M)

```

B-TREE-INSERT-NONFULL( $x = T.root$ , M)
  B-TREE-INSERT-NONFULL( $x = x.c_2$ , M)
    B-TREE-SPLIT-CHILD( $x, 2$ )
      B-TREE-INSERT-NONFULL( $x = x.c_3$ , M)

```

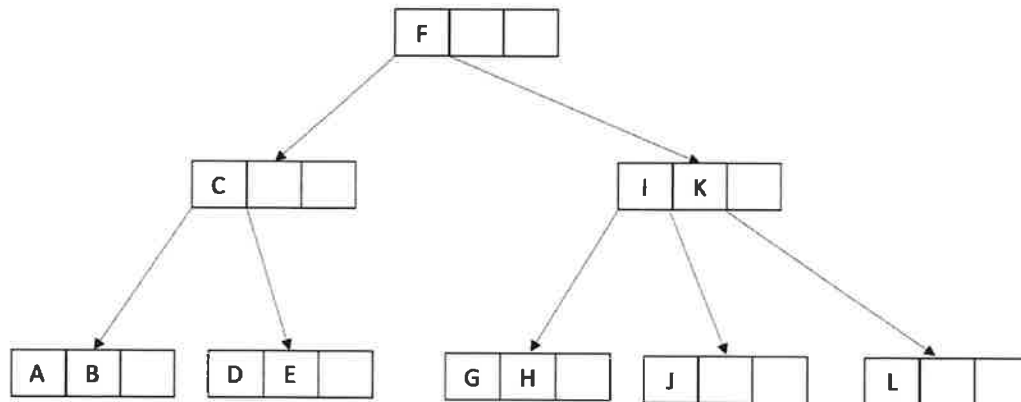
9/15

2. Visually represent the result of deleting from the BTree created in problem 1 the following values in the order listed:

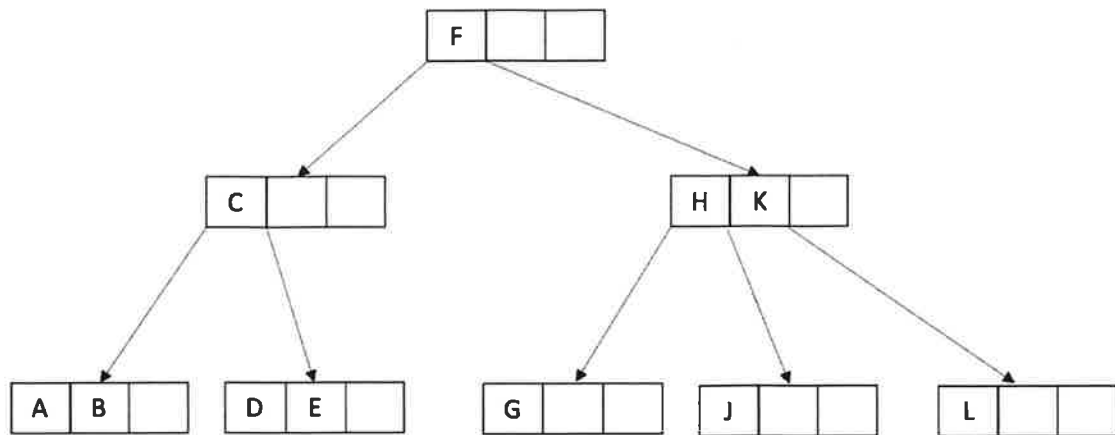
M, I, H, B, E

You are only responsible for showing the final tree after **each** deletion. No function calls are required.

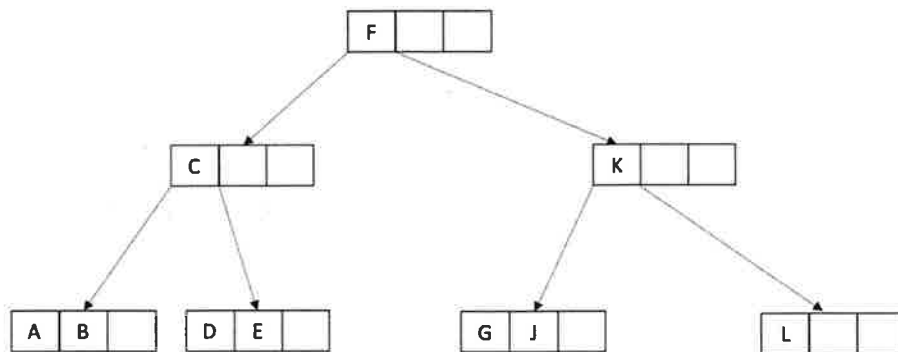
Delete M:



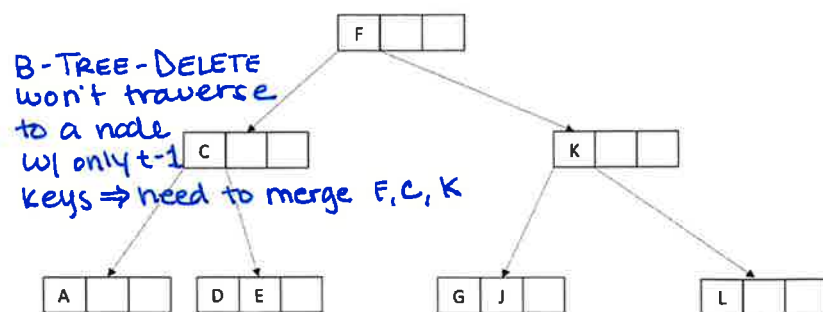
Delete I:



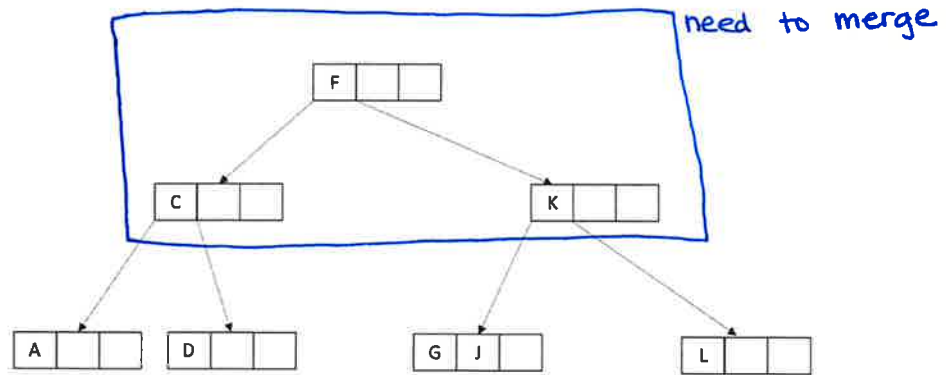
Delete H:



Delete B:



Delete E:



Pseudocode

- 4A/53 3. Write pseudocode for the public B-TREE class method B-TREE-DELETE(T, k).
 Pre-condition: The B-tree is valid with order t and the key to be deleted exists in the tree.
 Post-condition: The key is deleted from the B-tree, and the B-tree remains balanced.

Algorithm 1 Deleting a key k from a BTree T 12

```

procedure B-TREE-DELETE( $T, k$ )
  if  $T.root \neq NIL$  then
    B-TREE-DELETE( $T.root, k$ )
  end if
end procedure
  
```

when do we need to update $T.root$?

Pre-condition: The node x is valid and the key to be deleted exists in the tree.

Post-condition: The key is deleted from the B-tree rooted at node x , and the B-tree rooted at node x remains balanced.

Algorithm 2 Deleting a key k from a BTree rooted at a node x +5

```

procedure B-TREE-DELETE( $x, k$ )
  if  $x == NIL$  then
    return
  end if
   $i = \text{FIND-K}(x, k)$ 
  if  $i \leq x.n$  and  $x.keys_i == k$  then
    if  $x.leaf == \text{True}$  then
      REMOVE-LEAF-KEY( $x, i$ )
    else
      REMOVE-INTERNAL-KEY( $x, i, i + 1$ )
    end if
  else
    if  $x.leaf == \text{True}$  then
      return
    end if
     $y = x.c_i$ 
    if  $y.n == t - 1$  then
      if  $i > 1$  and  $x.c_{i-1}.n \geq t$  then
        SWAP-LEFT( $x, y, x.c_{i-1}, i - 1$ )
      else if  $i < x.n + 1$  and  $x.c_{i+1}.n \geq t$  then
        SWAP-RIGHT( $x, y, x.c_{i+1}, i$ )
      else if  $i > 1$  then
        MERGE-LEFT( $x.c_{i-1}, y, x.keys_i$ )
        DELETE( $x.keys_{i+1}$ )
      else
        MERGE-RIGHT( $x.c_{i+1}, y, x.keys_i$ )
        DELETE( $x.keys_i$ )
      end if
      B-TREE-DELETE( $y, k$ )
    end if
  end if
end procedure

```

not a node - need to move down all keys and children to account for x having 1 fewer child/key

Pre-condition: The node x is valid, and the k is the key to be searched in the node x .

Post-condition: If there exists an index i such that $k \leq x.keys_i$, then i is the smallest index for which this condition holds, else i is set to $x.n + 1$.

Algorithm 3 Determine the index i of the first key in a B-Tree node x where $k \leq x.key_i$
 ($i = x.n + 1$ if no such key exists) + 4

```

procedure FIND-K( $x, k$ )
  if  $x == NIL$  then
    return  $NIL$ 
  end if
   $i = 1$ 
  while  $i \leq x.n$  and  $k > x.keys_i$  do
     $i = i + 1$ 
  end while
  return  $i$ 
end procedure

```

Pre-condition: The node x is a valid leaf node and the index $i \leq x.n$.

Post-condition: The key at index i is deleted from the node x .

Algorithm 4 Remove the key at index i from a B-Tree leaf node x + 4

```

procedure REMOVE-LEAF-KEY( $x, i$ )
  if  $x == NIL$  or  $i > x.n$  then
    return  $x$ 
  end if
  for  $j = i; j < x.n; j = j + 1$  do
     $x.keys_j = x.keys_{j+1}$ 
  end for
  DELETE( $x.keys_j$ )
   $x.n = x.n - 1$ 
end procedure

```

Pre-condition: The node x is a valid node and $i \leq x.n$, $j \leq x.n + 1$.

Post-condition: The key at index i is deleted from the node x , the child at index j is modified accordingly, and the B-tree rooted at node x remains satisfied.

Algorithm 5 Remove the key at index i and child at index j from a B-Tree internal node x 4

procedure REMOVE-INTERNAL-KEY(x, i, j)

if $x == NIL$ or $i > x.n$ or $j > x.n + 1$ **then**

return x

end if

$k = x.keys_i$

$y = x.c_i$

$z = x.c_j$

if $y.n \geq t$ **then**

$pred = \text{MAX}(y)$

 B-TREE-DELETE($y, pred$)

$x.keys_i = pred$

else if $z.n \geq t$ **then**

$suc = \text{MIN}(z)$

 B-TREE-DELETE(z, suc)

$x.keys_i = suc$

else

 MERGE-LEFT(y, z, k)

for $l = i; l < x.n; l = l + 1$ **do**

$x.keys_l = x.keys_{l+1}$

end for

~~DELETE($x.keys_i$)~~

 DELETE(z)

 B-TREE-DELETE(y, k)

end if

$x.n = x.n - 1$ not necessarily if k was replaced by $pred$ or suc

end procedure

need to also shift children down
now that contents of z are in y

Pre-condition: The node x is a valid node.

Post-condition: The maximum key of the sub-tree rooted at x is returned and the tree is unchanged.

Algorithm 6 Return the maximum key in the B-Tree rooted at x +3

```

procedure MAX( $x$ )
  if  $x == NIL$  then
    return  $x$ 
  end if
  while  $x.leaf$  not  $True$  can directly use booleans :) do
     $x = x.c_{x.n+1}$ 
  end while
  return  $x.keys_{x.n}$ 
end procedure

```

Pre-condition: The node x is a valid node.

Post-condition: The minimum key of the sub-tree rooted at x is returned and the tree is unchanged.

Algorithm 7 Return the minimum key in the B-Tree rooted at x +3

```

procedure MIN( $x$ )
  if  $x == NIL$  then
    return  $x$ 
  end if
  while  $x.leaf$  not  $True$  can directly use booleans :) do
     $x = x.c_1$ 
  end while
  return  $x.keys_1$ 
end procedure

```

Pre-condition: The node x and its right sibling y are valid nodes that have $t - 1$ keys, and k is a valid key value.

Post-condition: The key k and node y is merged into x , and the B-tree properties are maintained.

Algorithm 8 Merge key k and all keys and children from y into y 's left sibling x + 6

```

procedure MERGE-LEFT( $x, y, k$ )
  if  $x == NIL$  or  $x.n \neq t - 1$  or  $y == NIL$  or  $y.n \neq t - 1$  then
    return
  end if
   $x.n = x.n + 1$ 
   $x.keys_{x.n} = k$ 
  for  $i = 1; i \leq y.n; i = i + 1$  do
     $x.keys_{x.n+i} = y.keys_i$ 
  end for
  for  $i = 1; i \leq y.n + 1; i = i + 1$  do
     $x.C_{x.n+i} = y.C_i$ 
  end for
   $x.n = 2t - 1$  fair, cleaner to use  $x.n += y.n$ 
end procedure

```

Pre-condition: The node x and its left sibling y are valid nodes that have $t - 1$ keys, and k is a valid key value.

Post-condition: The key k and node y is merged into x , and the B-tree properties are maintained.

Algorithm 9 Merge key k and all keys and children from y into y 's right sibling x + 4

```

procedure MERGE-RIGHT( $x, y, k$ )
  if  $x == NIL$  or  $x.n \neq t - 1$  or  $y == NIL$  or  $y.n \neq t - 1$  then
    return
  end if
   $x.n = x.n + 1$ 
   $x.keys_{x.n} = k$  should be smaller than all keys already in  $x$ 
  for  $i = 1; i < x.n; i = i + 1$  do
     $x.keys_{x.n+i} = x.keys_i$ 
  end for
  for  $i = 1; i \leq y.n; i = i + 1$  do
     $x.keys_i = y.keys_i$ 
  end for
  for  $i = 1; i \leq x.n; i = i + 1$  do
     $x.c_{x.n+i} = x.c_i$ 
  end for
  for  $i = 1; i \leq y.n + 1; i = i + 1$  do
     $x.c_i = y.c_i$ 
  end for
   $x.n = 2t - 1$  fair ☺
end procedure

```

Pre-condition: The node y , its left sibling z and its parent x are valid nodes. i is the index dividing children y and z in x . Node y is at least 1 key away from being full and Node z is at least 1 key away from being min-ed out.

Post-condition: Node y has an extra element that belonged to node x and Node x has an extra element that belonged to node z . B-tree property is maintained.

Algorithm 10 Give y an extra key by moving a key from its parent x down into y , moving a key from y 's left sibling z up into x , and moving the appropriate child pointer from z into y .
 Let i be the index of the key dividing y and z in x . 15

procedure SWAP-LEFT(x, y, z, i)

if $z.n < t - 2$ or $y.n > 2t - 2$ **then**

return

end if

$k = y.n$

while $k > 0$ **do**

$y.keys_{k+1} = y.keys_k$

$k = k - 1$

end while

if ~~$y.leaf$~~ ~~$== \text{False}$~~ **then**

$j = y.n + 1$

while $j > 0$ **do**

$y.c_{j+1} = y.c_j$

$j = j - 1$

end while

end if

$y.keys_1 = x.keys_i$

if ~~$y.leaf$~~ ~~$== \text{False}$~~ **then**

$y.c_0 = z.c_{z.n+1}$

end if

$x.keys_i = z.keys_{z.n+1}$

$y.n = y.n + 1$

$z.n = z.n - 1$

end procedure

careful. we index starting @ 1

Pre-condition: The node y , its right sibling z and its parent x are valid nodes. i is the index dividing children y and z in x . Node y is at least 1 key away from being full and Node z is at least 1 key away from being min-ed out.

Post-condition: Node y has an extra element that belonged to node x and Node x has an extra element that belonged to node z . B-tree property is maintained.

Algorithm 11 Give y an extra key by moving a key from its parent x down into y , moving a key from y 's right sibling z up into x , and moving the appropriate child pointer from z into y .

Let i be the index of the key dividing y and z in x . ✓ 4

procedure SWAP-RIGHT(x, y, z, i)

if $z.n < t - 2$ or $y.n > 2t - 2$ **then**

return

end if

$y.keys_{y.n} \leftarrow x.keys_i$

if ~~$y.leaf$~~ not ~~$False$~~ **then**

$x.c_{x.n+1} = z.c_1$ child should be added to y

end if

$x.keys_i = z.keys_0$

$k = 2$

while $k \leq z.n$ **do**

$z.keys_{k-1} = z.keys_k$

$k = k + 1$

end while

if ~~$z.leaf$~~ not ~~$False$~~ **then**

$j = 2$

while $j \leq z.n$ +1 **do**

$z.c_{j-1} = z.c_j$

$j = j + 1$

end while

end if

$y.n = y.n + 1$

$z.n = z.n - 1$

end procedure

6/6 **C++ Implementation** nice, thorough work here!

4. How would each pseudocode option impact the runtime of the B-TREE-INSERT function? Consider both asymptotic analysis as well as real-time impacts.

The main difference between the two ALLOCATE-NODE options is that the one that uses a vector implementation does not initialize the size of the vector while the one that uses an array does. Because an array is sized dynamically, this does not pose a problem with compilation, but does with runtime. This is because whenever we try to insert into a full vector, the vector would move to another space in memory and create a copy of itself there as a way to dynamically allocate more space. This results in a slower runtime, as shown below.

The function for B-Tree-Insert(T, k) is as follows:

```

B-TREE-INSERT( $T, k$ )
1   $r = T.root$ 
2  if  $r.n == 2t - 1$ 
3       $s = \text{ALLOCATE-NODE}()$ 
4       $T.root = s$ 
5       $s.leaf = \text{FALSE}$ 
6       $s.n = 0$ 
7       $s.c_1 = r$ 
8      B-TREE-SPLIT-CHILD( $s, 1$ )
9      B-TREE-INSERT-NONFULL( $s, k$ )
10 else B-TREE-INSERT-NONFULL( $r, k$ )

```

Lines 1-7 of B-Tree-Insert all take up constant time $O(1)$. The different implementations of Allocate-Node are both $O(1)$ time since assigning values to variables takes up $O(1)$ time and the vector initialized originally is NULL.

Lines 8-10 is where we can see the differences in implementation. The code for B-Tree-Split-Child($s, 1$) called in line 8 is as follows:

```

B-TREE-SPLIT-CHILD( $x, i$ )
1   $z = \text{ALLOCATE-NODE}()$ 
2   $y = x.c_i$ 
3   $z.leaf = y.leaf$ 
4   $z.n = t - 1$ 
5  for  $j = 1$  to  $t - 1$ 
6       $z.key_j = y.key_{j+t}$ 
7  if not  $y.leaf$ 
8      for  $j = 1$  to  $t$ 
9           $z.c_j = y.c_{j+t}$ 
10  $y.n = t - 1$ 
11 for  $j = x.n + 1$  downto  $i + 1$ 
12      $x.c_{j+1} = x.c_j$ 
13  $x.c_{i+1} = z$ 
14 for  $j = x.n$  downto  $i$ 
15      $x.key_{j+1} = x.key_j$ 
16  $x.key_i = y.key_t$ 
17  $x.n = x.n + 1$ 
18 DISK-WRITE( $y$ )
19 DISK-WRITE( $z$ )
20 DISK-WRITE( $x$ )

```

As stated above, when a vector overflows, it is copied into another memory location. Because the number of children of a B-tree Node can range from t to $2t$ and the number of keys in a B-tree Node can range from $t-1$ to $2t-1$, we can say that this copying operation takes $O(t)$ time for the implementation of **ALLOCATE-NODE** that uses a vector, where t is the min-degree of the B-tree. On the other hand, because

the array is initialized with size $2t-1$, inserting into an array will consistently take $O(1)$ time. This is reflected in lines 9, 12 and 15 of `B-TREE-SPLIT-CHILD`. If pushing a key or a child rightwards overflows the vector, it will take $O(t)$ time, while lines 9, 12 and 15 will only take $O(1)$ time if implemented with the array implementation.

`B-TREE-INSERT-NONFULL(x, k)`

```

1   $i = x.n$ 
2  if  $x.leaf$ 
3      while  $i \geq 1$  and  $k < x.key_i$ 
4           $x.key_{i+1} = x.key_i$ 
5           $i = i - 1$ 
6       $x.key_{i+1} = k$ 
7       $x.n = x.n + 1$ 
8      DISK-WRITE( $x$ )
9  else while  $i \geq 1$  and  $k < x.key_i$ 
10      $i = i - 1$ 
11      $i = i + 1$ 
12     DISK-READ( $x.c_i$ )
13     if  $x.c_i.n == 2t - 1$ 
14         B-TREE-SPLIT-CHILD( $x, i$ )
15         if  $k > x.key_i$ 
16              $i = i + 1$ 
17     B-TREE-INSERT-NONFULL( $x.c_i, k$ )

```

The while loop from lines 3-5 also follow this same logic. if $x.key_{i+1}$ is larger than the current size of the vector, this will overflow the vector and cause it to dynamically copy itself to a new memory location. This copying will take $O(t)$ time, while an array initialized first with the size will take $O(1)$ time to insert into.

Therefore, with both implementations, it will take $O(t)$ time asymptotically, but the multiple on t will be a lot larger if using the vector version of the implementation. Thus, while the two may seem the same in asymptotic analysis, the array implementation that initializes $2t-1$ sized arrays will be a lot faster in real life.